

- The JavaScript **event loop** and **run-to-completion** semantics.
- Callbacks for event handlers.
- Pyramid of doom.
- Taming asynchronous code: promises.
- Taming asynchronous code: `async`, `await`.

The Need for Concurrency

- Modern CPUs have clocks in the low GHz. That means individual CPU operations occupy under 1 nanosecond.
- Typically, I/O may take in the order of milliseconds which is around a million times slower than CPU operations.
- Highly inefficient to have CPU wait for an I/O operation to complete.
- Need to concurrently do other stuff while waiting for I/O to complete.
- Note that browser responsiveness is usually controlled by I/O responsiveness.

Approaches to Concurrency

Two commonly used approaches to concurrency:

Synchronous Blocking Model When a program attempts to do I/O, the program blocks until the I/O is completed. To allow concurrency, the operating system will schedule some other activity while waiting for the I/O. The unit of scheduling is usually a **process** or **thread**; leads to the process/thread model used by many current OS's.

Asynchronous Event Model When a program attempts to do I/O, it merely starts the I/O after registering a handler to handle the I/O completion event. The program continues running while the I/O is happening concurrently. The completion of the I/O results in an event which results in the registered handler being called.

JavaScript uses the asynchronous event model.

JavaScript Event Loop

The top-level JavaScript runtime consists of an event loop which pulls extant events off an event queue and calls their registered handlers:

```
while (eventQueue.notEmpty()) {  
    const event = eventQueue.remove();  
    const handler = event.handler();  
    handler.call(); //pass suitable arguments  
}  
//terminate program
```

- The `handler.call()` **runs to completion**.
- Code does not need to deal with an event handler being interrupted.
- Code still needs to deal with the fact that the order of running of event handlers is not defined.

Run to Completion Consequences

In `run-to-completion.js`:

```
#!/usr/bin/env nodejs
```

```
//BAD CODE!!
```

```
function sleep(seconds) {  
  const stop = Date.now() + seconds*1000;  
  while (Date.now() < stop) {  
    //busy waiting: yuck!  
  }  
}
```

```
setTimeout(() => console.log('timeout'),  
           1000 /*delay in milliseconds*/);
```

```
sleep(5);  
console.log('sleep done');
```

Run to Completion Log

```
12:02:17|master/code $ ./run-to-completion.js  
sleep done  
timeout  
12:02:23|master/code $
```

Because of run-to-completion semantics, it will **always** be the case that the *sleep done* message will be output before the *timeout* message.

Why This Concurrency Model

- JavaScript was designed as a language which should be easy for inexperienced programmers to use for scripting dynamic behavior in browsers.
- Browser reacts to user actions by generating events like key-press, mouse-click, etc.
- Browser programmer needs to provide optional event handlers for these events in order to implement browser dynamic behavior.
- Since every event handler runs to completion, programmer can simply concentrate on code for that event, ignoring other events (at least for independent events).
- No need for the programmer to understand complex process / threading models.
- Lower overhead for I/O bound tasks.

Playing with Asynchronous Functions

Use `later` to run function asynchronously after a random delay:

```
const MAX_TIMEOUT = 3;

function later(fn, ...args) {
  const timeout =
    Math.floor(Math.random()*(MAX_TIMEOUT + 1));
  setTimeout(fn, timeout*1000, ...args);
}
```

```
> .load later.js
undefined
> > later(() => console.log('done'))
undefined
> done //note prompt output before 'done'
```


Using Return Value of Asynchronous Function

```
> function f() {  
    later(() => { console.log('f run'); return 42; });  
}  
undefined  
> f()  
undefined  
> f run  
  
>
```

How do I get hold of the 42 return value.

Return Value of Asynchronous Function: Another Attempt

```
> let ret = -1
undefined
> function f() {
  later(() => { console.log('f run'); ret = 42; });
}
undefined
> f(); console.log('ret = ${ret}')
ret = -1
undefined
> f run

> ret
42
>
```

Passing a Handler for Return Value

```
> function f(succFn) {  
    later(() => { console.log('f run'); succFn(42); });  
}  
undefined  
> f((ret) => console.log('ret = ${ret}'))  
undefined  
> f run  
ret = 42  
  
>
```

Passing Return Value to Another Asynchronous Function

```
> function g(v, fn) {  
  later(() => { console.log('g(${v})'); fn(2*v); });  
}  
undefined  
> f((v) => g(v, (x) => console.log('g() value = ${x}')))  
undefined  
> f run  
g(42)  
g() value = 84  
  
>
```

Getting out-of-hand!

```
//Normal exception catching  
> try {  
    throw 'throwing';  
} catch (ex) {  
    console.log('caught ${ex}');  
}  
caught throwing  
undefined
```

Errors Continued

```
//Exception in Async not caught
> try {
    f(() => { throw { msg: 'thrown' }; })
  }
  catch (ex) {
    console.log('caught ${ex}');
  }
undefined
> f run
Thrown: [object Object]
>
```

Problems with Callbacks

- A top-level exception handler does not work for asynchronous callbacks since the handler runs before the callback. Hence exceptions occurring within the callback are not caught by the top-level exception handler.
- If a asynchronous function result needs to be further processed by another asynchronous function, then we need to have nested callbacks.
- A chain of callbacks leads to the **pyramid of doom** because of nesting of callbacks.

Promises

- A **Promise** is an object representing the eventual completion or failure of an asynchronous operation.
- When a function which requires an asynchronous callback as an argument is called, it returns immediately with an object called a *pending Promise*. Subsequently, the callbacks can be added to the promise. The callbacks will be called after the promise has been *settled*.

```
let promise = some_call_which_returns_promise(...);  
promise.  
  then(callback1).  
  then(callback2).  
  ...  
catch(errorCallback);
```


Promise Advantages

- Promises can be chained; this avoids the *pyramid of doom*.
- Callbacks are never called before completion of current run of js **event loop**.
- Callbacks added using `then` even after completion of the asynchronous operation will still be called.
- `then()` can be called multiple times to add multiple callbacks (called in order of insertion).
- Allows catching errors much more easily using `catch()`; similar to exception handling.
- `then()` can even be chained after a `catch`.

```
new Promise(  
  /* executor */  
  function(resolve, reject) { ... }  
);
```

- Creates a promise.
- `resolve` and `reject` are single argument functions.
- Executor function executed immediately. Usually will start some kind of asynchronous operation which may return some result.
 - 1 If the async operation succeeds with some result `value`, then the executor function should call `resolve(value)`.
 - 2 If the async operation fails with some error `err`, then the executor function should call `reject(err)`.

Promise States

Pending The underlying operation is not yet complete.

Fulfilled The underlying operation completed successfully.

Rejected The underlying operation failed.

Settled The operation completed, the promise is either *fulfilled* or *rejected*.

A promise is settled only once. The state of the promise will not change once it is settled.

Getting Promise Settlement: `then()`

```
somePromise.then(value, err)
```

- Arguments are one argument functions called when `somePromise` is settled; specifically `value` / `err` are called with fulfillment / rejection value depending on settlement.
- Usually `then()` is called with only the `value` argument, with rejection of `somePromise` handled using a `catch()`.
- `then()` itself returns a promise; this allows chaining `then`'s.
 - If the function passed to `then()` returns a value, then the return'd promise fulfills with that value.
 - If the function passed to `then()` throws an error, then the return'd promise rejects with that error.
 - If the function passed to `then()` returns a promise, then the return'd promise has the same settlement as it.

Getting Promise Rejection: `then()`

```
somePromise.catch(err)
```

- `err` is a one argument functions called with the rejection value of promise `somePromise`.
- `catch()` itself returns a promise; this allows continued promise chaining. Return value is similar to that of `then()`.

Playing with Promises

```
> function p(...args) { console.log(...args); }  
> p(1, 2)  
1 2  
> pr = new Promise((resolve, reject) => resolve(22))  
Promise { 22, ... }  
> pr.then((v) => p(v))  
Promise { <pending>, ... }  
> 22  
//Promise is settled only once  
> pr = new Promise((succ) => { succ(42); succ(22); })  
Promise { 42, ... }  
> pr.then((v) => p(v))  
Promise { <pending>, ... }  
> 42
```

Playing with Promises: Chaining then()'s

```
> function f(a, b) { p(a); return a * b; }  
undefined  
> pr = new Promise((resolve) => resolve(22))  
Promise { 22, ... }  
> pr.then((val) => f(val, 2)).  
    then((val) => f(val, 3)).  
    then((val) => p(val))  
> Promise { <pending>, ... }  
> 22  
44  
132  
  
>
```

Creating Settled Promises

`Promise.resolve(value)` Returns a promise which is already fulfilled with `value`.

`Promise.reject(err)` Returns a promise which is already rejected with `err`.

Playing with Promises: Asynchronous Functions

```
> function f(a, b, ret) {  
    p(`${new Date().toISOString()}: ${a}`);  
    setTimeout(() => ret(a*b), 2000);  
}
```

undefined

```
> pr = Promise.resolve(22)  
> pr.  
  then((v) => new Promise((succ) => f(v, 2, succ))).  
  then((v) => new Promise((succ) => f(v, 3, succ))).  
  then((v) =>  
    p(`${new Date().toISOString()}: ${v}`))
```

```
> 10:54:50 GMT-0500 (EST): 22
```

```
10:54:52 GMT-0500 (EST): 44
```

```
10:54:54 GMT-0500 (EST): 132
```

```
>
```

Playing with Promises: Errors

```
> function t() { return new Date().toString(); }
> pr1 = Promise.reject(new Error(t()))
Promise { <rejected> Error: 11:12:04 ... }
> (node:24159) UnhandledPromiseRejectionWarning: ...
> p(t()); pr1 =
    Promise.reject(new Error(t())); pr1.catch(()=>{})
11:15:36 GMT-0500 (EST)
...
> p(t()); pr1.
    then((v) => p(v)).
    then((v) => p(v)).catch((err)=>p(err))
11:16:10 GMT-0500 (EST)
...
> Error: 11:15:36 GMT-0500 (EST)
...
```

Playing with Promises: Errors Continued

```
> pr1.  
  then((v) => p('got value ${v}')).  
  then((v) => p('got value ${v}')).  
  catch((e) => { p('caught ${e}'); return 42; }).  
  then((v) => p('got value ${v}'))
```

```
Promise { <pending>, ... }
```

```
> caught Error: 11:15:36 GMT-0500 (EST)  
got value 42
```

```
>
```

Playing with Promises: Errors Continued

then()-chain continues past catch():

```
> Promise.resolve(1).  
  then((v) => { p('then1: ${v}'); return v*2; }).  
  then((v) => { p('then2: ${v}'); return v*2; }).  
  catch((e) => p('caught ${e}')).  
  then((v) => { p('then3: ${v}'); return v*2; })
```

```
Promise { <pending>, ... }
```

```
> then1: 1
```

```
then2: 2
```

```
then3: 4
```

```
>
```

Promise.all()

Given an **iterable** of promises, returns a promise containing array of fulfilled values, or rejection if any promise rejected.

```
> Promise.all([mul2(3), mul3(4), mul4(5)]).  
  then((v) => p(v))  
Promise { <pending>, ... }  
> [ 6, 12, 20 ]
```

```
> Promise.all([mul2(3), err(3)(2), mul3(4), mul4(5)]).  
  then((v) => p(v)).  
  catch((e) => p('caught ${e}'))  
Promise { <pending>, ... }  
> caught Error: err
```

Promise.all() Continued

Promise.all() runs all promises in parallel:

```
> p(t()); Promise.all([mul2(3), mul3(4), mul4(5)]).  
  then((v) => p(`${t()}: ${v}`))  
15:49:41 GMT-0500 (EST)  
Promise { <pending>, ... }  
> 15:49:43 GMT-0500 (EST): 6,12,20
```

Took 2 seconds to run all 3 functions even though each function takes 2 seconds apiece.

Promise.race()

Given an **iterable** of promises, returns a promise containing settlement of which ever incoming promise completes first.

```
> Promise.race([mul2(3), mul3(4), mul4(5)]).  
  then((v) => p(v))
```

```
Promise { <pending>, ... }
```

```
> 6
```

```
>
```

A Glimpse at Generators

Generators defined using `function*` and `yield`.

```
> function* seq(lo=0, hi=Number.POSITIVE_INFINITY) {  
    for (let i = Math.floor(lo); i <= hi; i++) yield(i);  
}
```

undefined

```
> for (s of seq(1, 3)) console.log(s);
```

1

2

3

undefined

```
>
```


Generators Return Iterators

- When a generator is called it does not run the generator code, but immediately returns an **iterator**.
- Generator code can `yield` successive values; `return` terminates the generator.
- Caller interacts with returned iterator to step the generator.
- Iterators have a `next()` method which returns an object with two properties:
 - `done` A boolean which is true when the generator is done.
 - `value` The currently yielded value.
- Passing argument to `next()` makes argument the value returned by `yield`.

Waiting For An Async Value Using Generators

```
function* asyncFn() {  
  const value = yield new Promise((resolve) => {  
    setTimeout(() => resolve(42), 2000);  
  });  
  //we can access value  
  console.log(value);  
}
```

```
//drive generator  
const iterator = asyncFn();  
const iteration = iterator.next();  
iteration.value.then(v => iterator.next(v));
```

async / await

- Converts promise code into synchronous style by yielding promises from generators.
- If a function or function expression has the `async` (contextual) keyword in front of it, then that function always returns a promise.
- When the `await` (contextual) keyword is used in front of an expression which is a promise, it blocks the program until the promise is settled. The value of an `await` expression is the fulfillment value of the promise.
- The `await` keyword can only be used within a `async` function.
- Errors can be handled using `try-catch`.
- Seems a big win.
- Note that we may need to fall back on promises using `Promise.all()` when we want to run code in parallel rather than sequentially.

async / await Example

```
> function msgPromise() {  
    return new Promise(function (resolve) {  
        setTimeout(() => resolve('hello@${t()}'),  
2000)})();  
}
```

undefined

```
> async function msg(n) {  
    const m = await msgPromise();  
    return `${n}: ${m}`  
}
```

undefined

async / await Example: Invoking using IIFE

```
> ( async () => {           //must use async to use await  
    p(await msg(22));  
    p(await msg(42));  
  })()                     //async IIFE
```

```
Promise { <pending>, ... }
```

```
> 22: hello@21:06:53 GMT-0500 (EST)
```

```
42: hello@21:06:55 GMT-0500 (EST)
```

```
>
```

Async sleep()

```
> async function sleep(millis) {  
  return new Promise((resolve) =>  
    setTimeout(() => resolve(), millis));  
}  
> (async function() {  
  p(t()); await sleep(2000); p(t()); }  
)()  
14:12:38 GMT-0500 (EST)  
Promise { <pending>, ... }  
> 14:12:40 GMT-0500 (EST)
```